

Horn clauses and negations

An **integrity constraint** is a clause of the form

$$\text{false} \text{ :- } a_1, \dots, a_k$$

where each a_i is an atom and `false` is a special atom that is false in all interpretations.

Horn clauses and negations

An **integrity constraint** is a clause of the form

$$\text{false} \text{ :- } a_1, \dots, a_k$$

where each a_i is an atom and **false** is a special atom that is false in all interpretations.

A **Horn clause** is either a definite clause or an integrity constraint.

Horn clauses and negations

An **integrity constraint** is a clause of the form

$$\text{false} :- a_1, \dots, a_k$$

where each a_i is an atom and **false** is a special atom that is false in all interpretations.

A **Horn clause** is either a definite clause or an integrity constraint.

The **negation** of a formula α , written $\neg\alpha$, is a formula that is true in an interpretation I iff α is false in I .

Horn clauses and negations

An **integrity constraint** is a clause of the form

$$\text{false} \text{ :- } a_1, \dots, a_k$$

where each a_i is an atom and **false** is a special atom that is false in all interpretations.

A **Horn clause** is either a definite clause or an integrity constraint.

The **negation** of a formula α , written $\neg\alpha$, is a formula that is true in an interpretation I iff α is false in I .

Example

$$KB = \begin{cases} \text{false} \text{ :- } a, b. \\ a \text{ :- } c. \\ b \text{ :- } c. \end{cases}$$
$$KB \models \neg c$$

Satisfiability

Every set of definite clauses is satisfiable (i.e., has a model).

Not so with Horn clauses

$$KB \models \varphi \iff KB, \neg\varphi \text{ is not satisfiable}$$

Satisfiability

Every set of definite clauses is satisfiable (i.e., has a model).

Not so with Horn clauses

$$\begin{aligned} KB \models \varphi &\iff KB, \neg\varphi \text{ is not satisfiable} \\ &\iff KB, \neg\varphi \models \text{false}. \end{aligned}$$

Satisfiability and disjunctions

Every set of definite clauses is satisfiable (i.e., has a model).

Not so with Horn clauses

$$\begin{aligned} KB \models \varphi &\iff KB, \neg\varphi \text{ is not satisfiable} \\ &\iff KB, \neg\varphi \models \text{false}. \end{aligned}$$

The **disjunction** $\alpha \vee \beta$ of α and β is a formula that is true in an interpretation I iff at least one of α or β is true in I .

Example

$$KB = \begin{cases} \text{false} :- a, b. \\ a :- c. \\ b :- d. \end{cases}$$
$$KB \models \neg c \vee \neg d$$

Satisfiability and disjunctions

Every set of definite clauses is satisfiable (i.e., has a model).

Not so with Horn clauses

$$\begin{aligned} KB \models \varphi &\iff KB, \neg\varphi \text{ is not satisfiable} \\ &\iff KB, \neg\varphi \models \text{false}. \end{aligned}$$

The **disjunction** $\alpha \vee \beta$ of α and β is a formula that is true in an interpretation I iff at least one of α or β is true in I .

Example

$$KB = \begin{cases} \text{false} :- a, b. \\ a :- c. \\ b :- d. \end{cases}$$
$$KB \models \neg c \vee \neg d$$

Horn-SAT is feasible, whereas 3-SAT is likely not.

Non-monotonicity

Logical consequence is **monotonic**: adding clauses doesn't invalidate a previous conclusion

$$KB \models \varphi \text{ implies } KB, \psi \models \varphi.$$

Non-monotonicity

Logical consequence is **monotonic**: adding clauses doesn't invalidate a previous conclusion

$$KB \models \varphi \text{ implies } KB, \psi \models \varphi.$$

Negation-as-failure leads to **non-monotonicity**: a conclusion can be invalidated by adding more clauses.

```
empty-course(X) :- course(X),  
                  \+enrolled(_,X).
```

Sometimes assume that a database of facts is complete.
Any fact not listed is false.

Non-monotonicity

Logical consequence is **monotonic**: adding clauses doesn't invalidate a previous conclusion

$$KB \models \varphi \text{ implies } KB, \psi \models \varphi.$$

Negation-as-failure leads to **non-monotonicity**: a conclusion can be invalidated by adding more clauses.

```
empty-course(X) :- course(X),  
                  \+enrolled(_,X).
```

Sometimes assume that a database of facts is complete.
Any fact not listed is false.

Example: Assume a database of video segments is complete.

Rules

Encode *birds fly*

```
fly(X) :- bird(X).
```

to allow for exceptions.

Rules and defaults

Encode *birds fly*

$\text{fly}(X) \text{ :- bird}(X).$ % $\frac{\text{bird}(X)}{\text{fly}(X)}$

to allow for exceptions.

Default rule (R. Reiter)

$$\frac{\text{bird}(X) : \text{fly}(X)}{\text{fly}(X)}$$

Rules and defaults

Encode *birds fly*

$\text{fly}(X) \text{ :- } \text{bird}(X).$ % $\frac{\text{bird}(X)}{\text{fly}(X)}$

to allow for exceptions.

Default rule (R. Reiter)

$$\frac{\text{bird}(X) : \text{fly}(X)}{\text{fly}(X)}$$

In general,

$$\frac{\text{prerequisite } p : \text{justification } j}{\text{conclusion } c}$$

applied to KB says:

conclude c if $KB \models p$ and $\underbrace{j \text{ is } KB\text{-consistent}}_{KB, j \not\models \text{false}}$

Rules and defaults

Encode *birds fly*

$\text{fly}(X) \text{ :- bird}(X).$ % $\frac{\text{bird}(X)}{\text{fly}(X)}$

to allow for exceptions.

Default rule (R. Reiter)

$$\frac{\text{bird}(X) : \text{fly}(X)}{\text{fly}(X)}$$

In general,

$$\frac{\text{prerequisite } p : \text{justification } j}{\text{conclusion } c}$$

applied to KB says:

conclude c if $KB \models p$ and $\underbrace{j \text{ is } KB\text{-consistent}}_{KB, j \not\models \text{false}}$

j is true in some model of KB

Birds and bees

$$\frac{\text{bird}(X) : \text{fly}(X)}{\text{fly}(X)}$$

Let KB be

`bird(robin).`

`bird(penguin).`

`false :- fly(penguin).`

`fly(bee).`

Conclude:

Birds and bees

$$(\star) \frac{\text{bird}(X) : \text{fly}(X)}{\text{fly}(X)}$$

Let *KB* be

```
bird(robin).  
bird(penguin).  
false :- fly(penguin).  
fly(bee).
```

Conclude:

fly(robin) by default rule (\star)

but *not* fly(penguin).

Birds and bees

$$(\star) \frac{\text{bird}(X) : \text{fly}(X)}{\text{fly}(X)}$$

Let KB be

```
bird(robin).  
bird(penguin).  
false :- fly(penguin).  
fly(bee).
```

Conclude:

`fly(robin)` by default rule (\star)

but *not* `fly(penguin)`.

An explanation of `fly(bee)` using (\star) is

`bird(bee)`

Birds and bees

$$(\star) \quad \frac{\text{bird}(X) : \text{fly}(X)}{\text{fly}(X)}$$

Let KB be

```
bird(robin).  
bird(penguin).  
false :- fly(penguin).  
fly(bee).
```

Conclude:

```
fly(robin)    by default rule (★)
```

but *not* fly(penguin).

An explanation of fly(bee) using (★) is

```
bird(bee)
```

which we can block by adding to KB the rule

```
false :- bird(bee).
```

Non-determinism

Conflicting defaults

$$\frac{\text{quaker}(X) : \text{pacifist}(X)}{\text{pacifist}(X)}$$

$$\frac{\text{republican}(X) : \text{hawk}(X)}{\text{hawk}(X)}$$

Non-determinism

Conflicting defaults

$$\frac{\text{quaker}(X) : \text{pacifist}(X)}{\text{pacifist}(X)}$$

$$\frac{\text{republican}(X) : \text{hawk}(X)}{\text{hawk}(X)}$$

Let KB be

`quaker(nixon).`

`republican(nixon).`

`false :- pacifist(X), hawk(X).`

Non-determinism

Conflicting defaults

$$\frac{\text{quaker}(X) : \text{pacifist}(X)}{\text{pacifist}(X)}$$

$$\frac{\text{republican}(X) : \text{hawk}(X)}{\text{hawk}(X)}$$

Let KB be

`quaker(nixon).`

`republican(nixon).`

`false :- pacifist(X), hawk(X).`

Applying one default to Nixon makes the other inapplicable.

KB has two incompatible extensions, breaking

least fixed point (provability model) for Horn clauses.

Normal default rules and inferring negations

A default rule is *normal* if its justification is its conclusion $\frac{p : c}{c}$
— infer c if it is consistent and p is provable

Normal default rules and inferring negations

A default rule is *normal* if its justification is its conclusion $\frac{p : c}{c}$
— infer c if it is consistent and p is provable

Closed World Assumption: any unprovable atom φ is false

$$\frac{\text{true} : \neg\varphi}{\neg\varphi}$$

Normal default rules and inferring negations

A default rule is *normal* if its justification is its conclusion $\frac{p : c}{c}$
— infer c if it is consistent and p is provable

Closed World Assumption: any unprovable atom φ is false

$$\frac{\text{true} : \neg\varphi}{\neg\varphi}$$

Negation as failure: φ is false if attempting to prove φ fails finitely

`naf(P) :- (P,!,fail); true.`

Normal default rules and inferring negations

A default rule is *normal* if its justification is its conclusion $\frac{p : c}{c}$
— infer c if it is consistent and p is provable

Closed World Assumption: any unprovable atom φ is false

$$\frac{\text{true} : \neg\varphi}{\neg\varphi}$$

Negation as failure: φ is false if attempting to prove φ fails finitely

`naf(P) :- (P,!,fail); true.`

N.B. Checking finite failure can be as hard as the Halting Problem.

3 modes of inference (C.S. Peirce)

		typed functional prog \cong proof
Deduction	deduce	modus ponens \cong function app $f(a)$
Abduction	explain	choose input a from <i>assumables</i>
Induction	generalise/program	choose rule/function f

3 modes of inference (C.S. Peirce)

		typed functional prog \cong proof
Deduction	deduce	modus ponens \cong function app $f(a)$
Abduction	explain	choose input a from <i>assumables</i>
Induction	generalise/program	choose rule/function f

From \models as inclusion \subseteq

$$\begin{aligned} KB \models g &\iff \text{Mod}(KB) \subseteq \text{Mod}(g) \\ KB \text{ satisfiable} &\iff \text{Mod}(KB) \not\subseteq \text{Mod}(\text{false}) \\ &\iff \text{Mod}(KB) \neq \emptyset \end{aligned}$$

to weighing alternatives $d \in D$ via probabilities given KB

$$\text{prob}(d|KB) = \text{conditional probability of } d \text{ given } KB$$

\rightsquigarrow Bayesian networks ...