



**Trinity College Dublin**

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

**Week 1**

# Parallel programming in Haskell

**CS4012**

Topics in Functional Programming

Glenn Strong <[Glenn.Strong@tcd.ie](mailto:Glenn.Strong@tcd.ie)>

# Parallel Programming

## Introduction

---

- **Parallelism as an idea is very cool, but it can also be very hard to work with.**
- **Parallelism has been talked about as a good fit for Functional Programming since the earliest days of FP languages.**
- **Immutable data and side-effect free computation combined has a lot of promise.**

# Parallel Programming



---

- The various compute cores need to coordinate their activities, meaning that we (the programmers) have to think about:
  - Race conditions
  - Data dependencies
  - Locking (and thus locking *problems*, like deadlocks)
- The big problem seems to be *shared mutable data*

# Parallel Programming

- The big problem seems to be *shared mutable data*.

When things execute in sequence it's easy to know where you are




Current instruction



# Parallel Programming

- But in parallel...



Current instruction



# Parallel Programming

---

- **Some terminology that we have to get right from the outset:**
  - *Parallelism* is using multiple compute elements (e.g. more than one CPU core) to perform a computation.
  - *Concurrency* is the use of more than one thread of program control at a time (which may, or may not, involve more than one processor).

# Parallel Programming

## An example: Fibonacci numbers

---

A small example to help us explore...

```
fib :: Integer → Integer
fib 0 = 0
fib 1 = 1
fib n = fib (n-1) + fib (n-2)
main = print $ fib 37
```

# Parallel Programming

## Compiler flags

GHC has some built-in options to enable:

- Multicore
- Performance analysis (profiling)
- Configuring the runtime system

```
ghc -threaded -rtsopts -eventlog ex1.hs
```

Line with the threaded runtime

enable passing options  
to the runtime

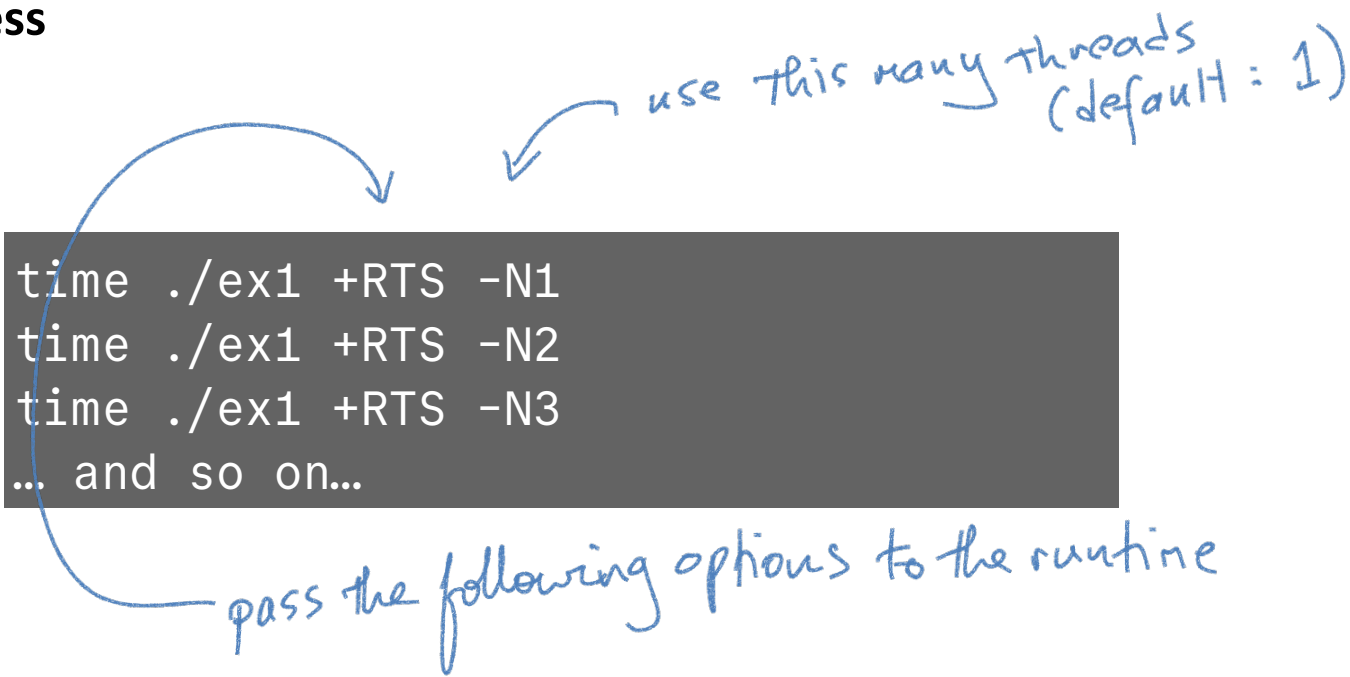
link a runtime that  
can record traces



# Parallel Programming

## Runtime flags

- Running it a few times with different options to the Run Time System will give us some data about how multiple cores speed up the process



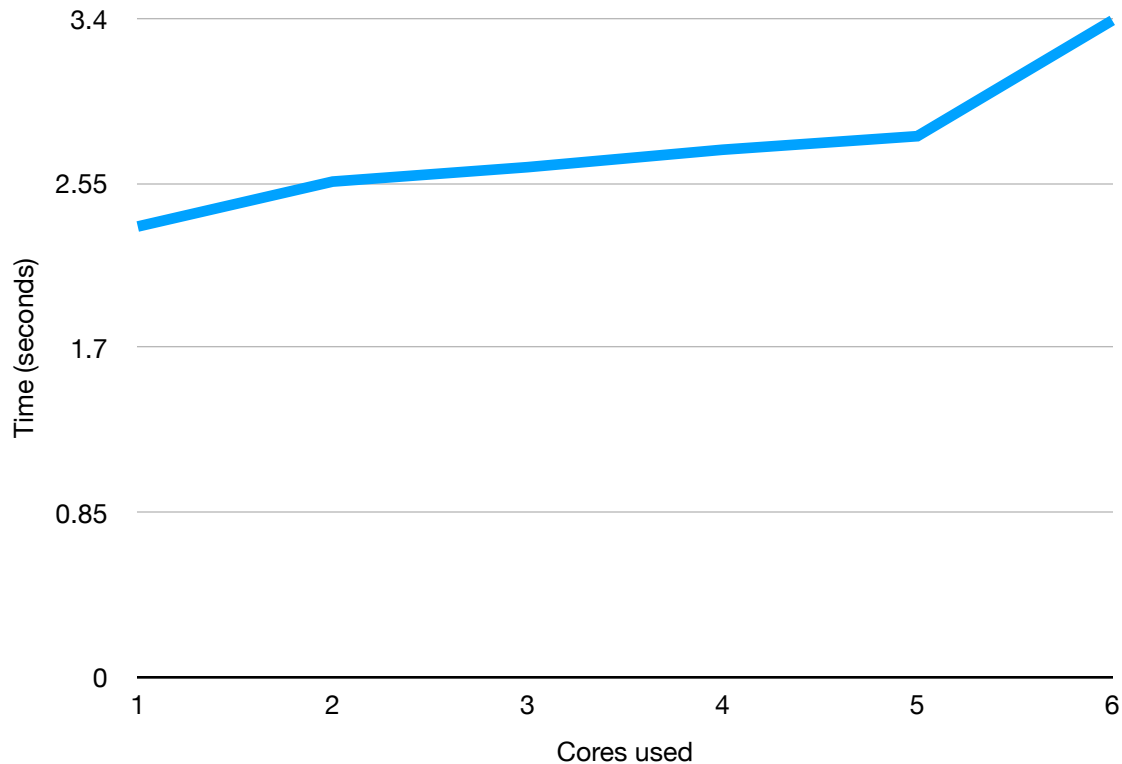
```
time ./ex1 +RTS -N1
time ./ex1 +RTS -N2
time ./ex1 +RTS -N3
... and so on...
```

use this many threads  
(default: 1)

pass the following options to the runtime

# Parallel Programming

## An example: Fibonacci numbers. Timings



# Parallel Programming

---

- This is very discouraging.
- The program did not get any faster when using multiple cores
  - In fact, it got a little bit slower!
- To find out what has gone wrong we will use the ThreadScope profiling tool.
- If you pass the “-l” flag to the runtime it will dump a log of various interesting events to a file (“ex1.eventlog”).
- This binary file is not the easiest thing to read on it’s own, we need

# Parallel Programming

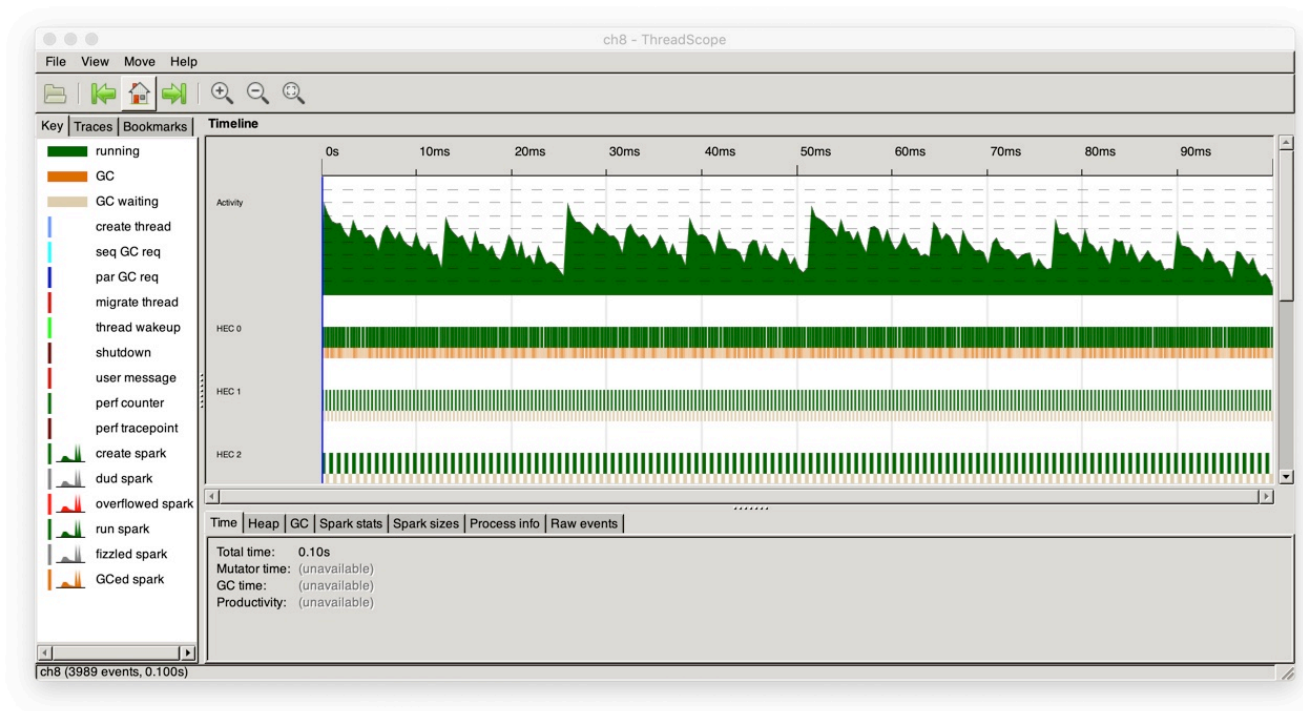
## ThreadScope

---

- ThreadScope is a profiling and analysis tool for parallel Haskell.
- It can be a little bit awkward to install from source
- I advise using a pre-built binary if you can. Instructions are on the Haskell Wiki and linked on Blackboard.
  - You need to install the Haskell GTK bindings
  - Then the ThreadScope binary
  - I will make a VM available for this if you need it.


# Parallel Programming

- The Threadscope UI



# Parallel Programming

- What do the ThreadScope results look like for our 'fib' program?





**Trinity College Dublin**  
Coláiste na Tríonóide, Baile Átha Cliath  
The University of Dublin

# Thank you

**[Glenn.Strong@scss.tcd.ie](mailto:Glenn.Strong@scss.tcd.ie)**

**<https://scss.tcd.ie/Glenn.Strong/>**



**Trinity College Dublin**

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

**Week 1**

# Parallel programming in Haskell

**CS4012**

Topics in Functional Programming

Glenn Strong <[glenn.Strong@scss.tcd.ie](mailto:glenn.Strong@scss.tcd.ie)>



# Parallel Programming

- The `Control.Parallel` library has a number of utilities to let us signal to the runtime that there are sites of potential parallelism.

```
par :: a → b → b
```

- It doesn't take long to persuade yourself that the only possible function that could have this type must be something equivalent to this:

```
par left right = right
```

- So what's the point?

# Parallel Programming

- We need to take a little digression. We know Haskell evaluates expressions *lazily*, but what does that mean.
- Here are two Haskell expressions entered at the GHCi prompt. They are very similar, but they are definitely different.

```
Prelude> let x = 1 + 2 :: Int  
Prelude> 3
```

- The difference is all to do with evaluation.
- We can't see the difference in normal use:


```
Prelude> x == 3  
True
```

# Parallel Programming

- But there is a difference.
- “x” could be represented as “3”
- or as a lazily-unevaluated computation of “1 + 2”.



Versus



# Parallel Programming

- **GHCi has a debugging operation to allow you to see whether something is evaluated or not:**

```
Prelude> let x = 1 + 2 :: Int
Prelude> :sprint x
x = _
Prelude> x
3
Prelude> :sprint x
x = 3
```


- **This unevaluated expression that has been printed as “\_” is sometimes called a “Thunk”.**

# Parallel Programming

- More complex examples could include the unevaluated “graphs” formed by something like this:


```
let x = 1 + 2  
let y = x + x
```

- What does all this have to do with parallelism?



# Parallel Programming

- Well, imagine if we have some expression graph with shared nodes.
- If the right sub tree was evaluated in parallel with the left sub tree then the result would be shared.



# Parallel Programming

---

- So, coming back to our function:

```
par :: a → b → b
```

- Semantically the expression `par x y` is equivalent to just `y`, but the runtime is allowed to use it as a hint that it would be a good idea to evaluate `x` in parallel with `y`

# Parallel Programming


---

Let's try again. This is ex2.hs:

```
fib :: Integer → Integer
fib 0 = 0
fib 1 = 1
fib n = par nf ( fib (n-1) + nf )
        where nf = fib (n-2)
```




# Parallel Programming




# Parallel Programming

- Threadscope actually looks OK at first...



# Parallel Programming

- Zooming in shows the problem. We get bursts of activity followed by complete stalls.



# Parallel Programming

```
fib :: Integer → Integer
fib 0 = 0
fib 1 = 1
fib n = par nf ( fib (n-1) + nf )
        where nf = fib (n-2)
```

**What's happening?**

**A new lazy task is started for `nf`. Nothing is demanding its evaluation, though...**

**The order that `(+)` evaluates its arguments is at the heart of this. It is *strict* in its left argument, so that forces `fib (n-1)` to be evaluated before the value of `nf` is demanded.**

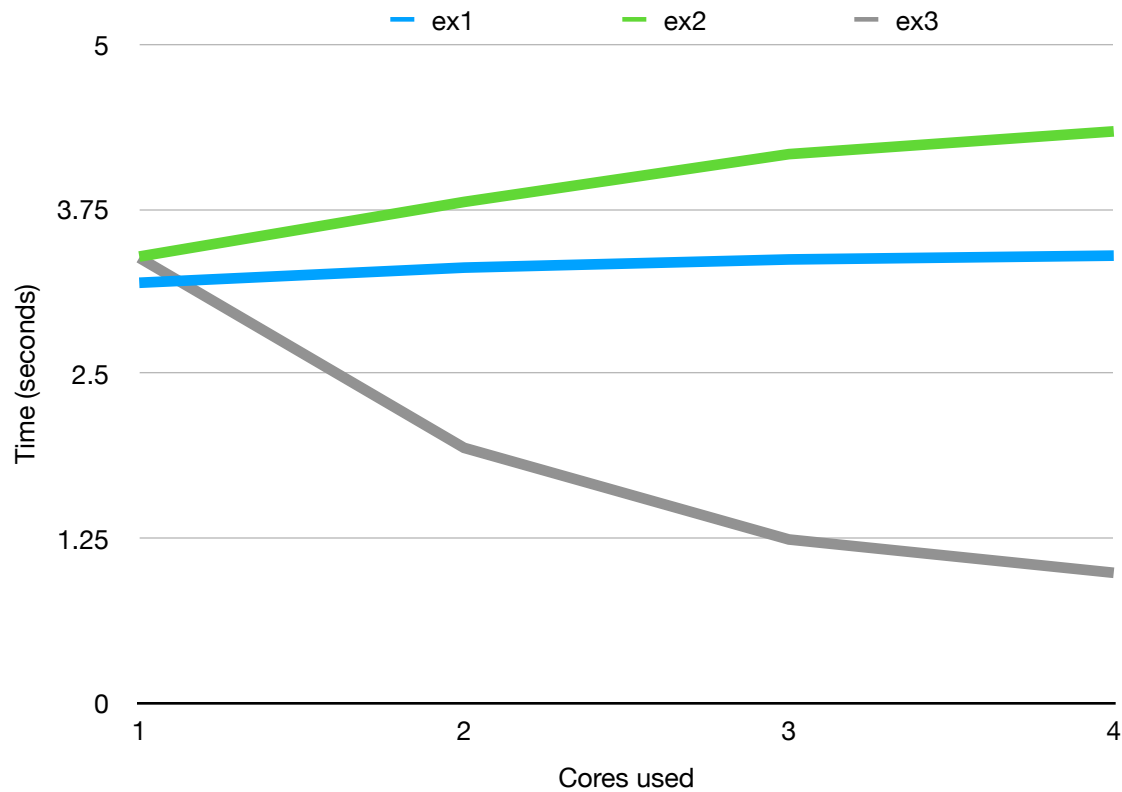
# Parallel Programming

So what would happen if we swapped the order of the arguments around?

```
fib :: Integer → Integer
fib 0 = 0
fib 1 = 1
fib n = par nf ( nf + fib (n-2) )
        where nf = fib (n-1)
```

**ex3.hs**

# Parallel Programming



# Parallel Programming

---

We shouldn't have to have this insider knowledge of how (+) treats its arguments, that was way too hard.

Introducing...

```
pseq :: a → b → b
```

pseq is like par but it is strict in its first argument.

# Parallel Programming

**pseq causes the main task to get on with nf2**

```
fib :: Integer → Integer
fib 0 = 0
fib 1 = 1
fib n = par nf1 (pseq nf2 (nf1 + nf2))
        where nf1 = fib (n-1)
              nf2 = fib (n-2)
```



# Parallel Programming

A sidebar on syntax. In Haskell we can write any *binary* function either like this:

```
f x y
```

or like this:

```
x `f` y
```

Sometimes it's easier to read the program using the infix notation, and so I will use it from time to time:

```
fib :: Integer → Integer
fib 0 = 0
fib 1 = 1
fib n = nf1 `par` nf2 `pseq` nf1 + nf2
      where nf1 = fib (n-1)
            nf2 = fib (n-2)
```

# Parallel Programming

Spark overhead can dominate after a while. If we limit new threads to allow a better distribution of work then we can get even better performance:

```
import Control.Parallel

sfib :: Integer → Integer
sfib n | n < 2 = 1
sfib n = sib (n-1) + sib (n-2)

fib :: Integer → Integer → Integer
fib 0 n = sfib n
fib _ n | n < 2 = 1
fib d n = nf1 `par` nf2 `pseq` nf2 (nf1 + nf2)
          where nf1 = fib (d-1) (n-1)
                nf2 = fib (d-1) (n-2)

main = print $ fib 3 37
```

# Parallel Programming

---

- You might think this is getting a little tricky to use.
- Lots of things to think about
  - Unevaluated vs Evaluated computation
  - Relative costs and sizes of computation
  - Sharing
- To explore how Haskell addresses these we need to bring in a big idea (maybe *the* big idea of this module...)



**Trinity College Dublin**

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

# Thank you

**Glenn.Strong@scss.tcd.ie**

**<https://scss.tcd.ie/Glenn.Strong/>**